# The BLASBench Report

Philip J. Mucci

Kevin S. London

John Thurman

mucci@cs.utk.edu

london@cs.utk.edu

thurman@cs.utk.edu

February, 1999

## 1   Introduction

BLASBench is a benchmark designed to evaluate the performance of some kernel operations of different implementations of the BLAS routines. The BLAS are the Basic Linear Algebra Subroutines and are found in some form or another on most vendors' machines. The BLAS were initially developed as part of LAPACK. Their goal was to provide a standardized API for common Vector-Vector, Vector-Matrix and Matrix-Matrix operations. A version of the BLAS is available from Netlib at `http://www.netlib.org`. This version is subsequently referred to as the reference version. The reference BLAS are completely unoptimized Fortran codes intended as a reference for correctness to the vendors.

As of November, 1998 BLASBench has been integrated into the LLCbench benchmarking suite.

## 2   Goals of BLASBench

BLASBench aims to provide the following:

- Evaluate the performance of the BLAS routines in MFLOPS/sec.

- Provide information for performance modeling of applications that make heavy use of the BLAS.

- Evaluate compiler efficiency by comparing performance of the reference BLAS versus the hand-tuned BLAS provided by the vendor.

- Validate vendor claims about the numerical performance of their processor.

- Compare against peak cache performance to establish bottleneck, memory or CPU.

# 3    Description

BLASBench currently benchmarks the three most common routines in the BLAS. They are:

- *AXPY* - Vector addition with scale

- *GEMV* - Matrix-vector multiplication with scale

- *GEMM* - Matrix-Matrix multiplication with scale

The benchmark can run in either single or double precision. This is important as many systems cannot sustain the additional memory bandwidth required by using double precision data. The test space is highly tunable, as are some of the metrics BLASBench reports. It reports its results in MFLOPS/sec and MB/sec. These numbers are not computed from hardware statistics, but rather from the absolute operation and memory reference count required by each algorithm.

# 4    How it works

BLASBench is a C program that calls BLAS routines written in Fortran, the reason for this being that we wish to perform dynamic memory allocation. For each test, BLASBench allocates its memory in such a way that the total amount of memory taken up by each test is less than or equal to the nearest power of two. The rationale for this is that our cache sizes are always in a power of two. Once the memory is allocated, the arrays are initialized, and the test is run for a certain number of iterations. By default, the iteration count is not constant over each problem size. BLASBench trys to keep the amount of data touched by each run constant. This means that larger problem sizes run for fewer iterations. The effect of this is that the run time for each size is approximately constant. BLASBench provides an option to keep the iteration count constant across all tests. After each size is tested, the cache is flushed and BLASBench either proceeds to the next size or repeats that size, depending on the options given to the program. BLASBench allows you to repeat each

size any number of times. This could be used to validate the numbers from each size on a time-shared system. By default, BLASBench calls the BLAS routines with the leading dimension of each array or vector set to the exact size of that vector. This means that the BLAS routines operate on the entire data set. Frequently however, BLAS routines are called upon smaller portions of larger arrays and matrices, thus an option is provided to keep the leading dimension constant among every test. In this case, BLASBench allocates the largest possible data set, and simply changes the working set size passed to each BLAS routine.

# 5   Using BLASBench

## 5.1   Obtain the Distribution

BLASBench is now found in the LLCbench distribution. The latest release of LLCbench can always be found through the original author's homepage at
http://icl.cs.utk.edu/projects/llcbench.

Now unpack the installation using gzip and tar.

```
kiwi> gzip -dc llcbench.tar.gz | tar xvf -
kiwi> cd llcbench
kiwi>  ls
Makefile     cachebench/  index.html   mpbench/     sys.def@
blasbench/   conf/        user.def     results
```

## 5.2   Build the Distribution

First we must configure the build for our machine, OS and BLAS libraries. All configurations support the reference BLAS if available. Before configuration make with no arguments lists the possible targets.

```
kiwi> make
Please use one of the following targets:

        solaris sunos5
        sun sunos4
        sgi-o2k o2k
        linux-mpich
```

```
        linux-lam
        alpha
        t3e
        ppc ibm-ppc
        pow2 ibm-pow2
        reconfig (to bring this menu up again)

After configuration, please check the VBLASLIB variable in
sys.def and make sure that it is pointing to the vendor BLAS
library if one exists.
```

Configure the build. Here, we are on a Solaris workstation.

```
kiwi> make solaris
ln -s conf/sys.solaris sys.def
```

BLASBench's default runtime variable values are contained in the file make.def and may be modified there.

Examine the sys.def file to ensure proper compiler flags and paths to the different BLAS libraries. The BLASLIB variable should contain the absolute path to the reference BLAS library and the VBLASLIB variable should contain the absolute path to the vendor's BLAS library. If one or the other is not available, just leave it blank and that specific executable will not be generated.

Now type make to get options for building a benchmark.

```
kiwi> make
Please use one of the following targets:

For all three  : bench, run, graphs
For Blasbench  : blas-bench, blas-run, blas-graphs
For Cachebench : cache-bench, cache-run, cache-graphs
For MPbench    : mp-bench, mp-run mp-graphs
```

Now build BLASBench. Depending on whether or not both BLASLIB and VBLASLIB are set, one or two executables will be generated.

```
kiwi> make blas-bench
cd blasbench; make blasbench; make vblasbench
/opt/SUNWspro/bin//cc -fast -dalign -DREGISTER -xarch=v8plusa -c bb.c
if [ -f "/src/icl/LAPACK_LIBS/blas_SUN4SOL2.a" ]; then /opt/SUNWspro/bin//f77 -xarch=v8p
-o blasbench bb.o /src/icl/LAPACK_LIBS/blas_SUN4SOL2.a -xlic_lib=sunperf; fi;
/opt/SUNWspro/bin//f77 -xarch=v8plusa -o vblasbench bb.o  -xlic_lib=sunperf
```

## 5.3   Running BLASBench

BLASBench can be run by hand, but it is intended to be run through the makefile. Running it via the makefile automates the collection and presentation process. By default, the makefile runs both executables with the arguments -c -o -e 1 -i 10 -x 2 -m 24. This says that the iteration count should be constant, the output should be reported in MFLOPS/sec, each size should be repeated only once, the iteration count should be set to ten, two measurements are taken betwen every problem size value that is a power of two, and the maximum problem size tested is $2^{24}$ bytes. You can change the default settings by changing the variables in make.def after you have configured the distribution.

```
kiwi> make blas-run
cd blasbench; make run
mkdir results
if [ -x blasbench ]; then blasbench  -i 10 -e 1 -m 24 -x 2 -c -d -o -v > results/daxpy.d
if [ -x blasbench ]; then blasbench  -i 10 -e 1 -m 24 -x 2 -c -d -o -a > results/dgemv.d
if [ -x blasbench ]; then blasbench  -i 10 -e 1 -m 24 -x 2 -c -d -o -t > results/dgemm.d
.
.
.
Now do a 'make blas-graphs'.
```

/em make blas-graphs will package up the datafiles for analysis on another machine. If Gnuplot is present, it will also generate the graphs in place and package them separately.

```
kiwi>make graphs
Z='hostname'; cd results; for i in saxpy daxpy sgemv dgemv sgemm dgemm; do if [ -r $.dat
Z='hostname'; cd results; for i in vsaxpy vdaxpy vsgemv vdgemv vsgemm vdgemm; do if  -r
Z='hostname'; cd results; mv info.dat $Z.info
```

```
Z='hostname'; Y='uname -m'; X='uname -sn'; sed -e "s|TITLE|UTK BLAS performance for Z: $
 .
 .
 .
Z='hostname'; cd results; tar cvf $Z-bp-datafiles.tar *.dat *.gp *.info;
 .
 .
 .
Z='hostname'; cd results; tar cvf $Z-bp-graphs.tar *.ps
blasperf.ps
compare.ps
vblasperf.ps

If you don't have GNUplot, you can make the graphs on another machine
using the blasbench/results/kiwi-bp-datafiles.tar file.
```

Using Gnuplot to create the graphs will result in either 1 or 3 graphs. Each graph contains the performance in megaflops of all three operations. They are named as follows:

- blasperf.ps - Postscript file of the reference BLAS.

- vblasperf.ps - Postscript file of the vendor's BLAS.

- compare.ps - Comparison of the two.

## 5.4  Arguments to BLASBench

This is the BLASBench arguement list from the command line help. The defaults listed are for direct execution of the benchmark (not the defaults for execution through the makefile).

```
kiwi> blasbench -h
Usage: blasbench [-vatsco -x # -m # -e # -i #]
        -v AXPY dot product benchmark
        -a GEMV matrix-vector multiply benchmark
        -t GEMM matrix-matrix multiply benchmark
        -s Use single precision floating point data
        -c Use constant number of iterations
```

6

```
-o Report Mflops/sec instead of MB/sec
-x Number of measurements between powers of 2.
-m Specify the log2(available physical memory)
-e Repeat count per cache size
-l Hold LDA and loop over sizes of square submatrices
-d Report dimension statistics instead of bytes
-i Maximum iteration count at smallest cache size


Default datatype   : double, 8 bytes
Default datatype   : float, 4 bytes
Defaults if to tty : -vat -x1 -m24 -e2 -i100000
Defaults if to file: -t   -x1 -m24 -e1 -i100000
```

# 6    Results on the CEWES MSRC Machines

The following graphs are taken from our runs on each of the CEWES machines during
dedicated time. Those machines are the SGI Origin 2000, the IBM SP and the Cray T3E.
The cache size and theoretical peak MFLOPS for each machine is listed as follows. The peak
MFLOPS is as reported by the vendor and is simply computed as a product of the clock
speed times the number of independant FMA's that can be computed per cycle.

| Machine | Cache | Peak |
|---------|-------|------|
| SGI Origin 2000 | 32K,4MB | 390 |
| IBM SP | 128K | 240 |
| Cray T3E | 8K,96K | 900 |

*It should be stated that AXPY and GEMV had a bug that was revealed during the writing
of this report. The bug was in the cache flushing code and the result was that part of the
operands to the BLAS still resided in cache, falsely inflating the numbers. The two graphs
for DAXPY and DGEMV are included at the end of this document for reference, however
the performance they indicate is erroneous.*

## 6.1    DGEMM

Figure 1: Performance of Matrix-Matrix Multiplication with Scale

Matrix multiply provides a lot of opportunity for cache reuse. By tiling the matrices, the working set can be reduced to the size of cache and thus only *capacity* misses are taken. For this reason, the performance of GEMM has long since served as a good indicator of a machine's *peak practical performance*. A machine with an adequate memory subsystem like the SP, can achieve very high efficiencies, that is high percentage of the vendor's published MFLOP rating.

The reader should notice that clock speed and L2 cache size do not play as critical a role in the performance of this routine as one might think. The spikes in the T3E's performance curve are due to the matrix dimension being a multiple of the block size. For other cases, the GEMM routine must engage in rather lengthy cleanup code. The small cache/line size of the T3E simply exaggerate the performance loss. The SP, with it's large line size and ability to issue 2 multiply-add instructions as well as a load/store per cycle does quite well, reaching approximately 90 percent efficiency. The Origin appears to suffer from it's small cache line size and its inability to issue a load/store every cycle. The R10000 processor, like the Power2, can also issue 2 multiply-adds per cycle, but it appears to have great difficulty keeping those functional units busy.

# 7 Future work

- Provide an option for measuring specific problem sizes and ranges.

- Provide an option to specify the problem sizes in dimensionality.

- Provide an option to specify the starting problem size.

- Use specialized, high-resolution timers where available.

- Add additional BLAS routines `TRSM`, `TRSV`, and `SYR2K`.

- Add parameters to tune the placement and padding of the arrays.

# 8 References

John Hennessey and David Patterson, *Computer Architecture, A Quantitative Approach*
    http://www.netlib.org/blas

## 8.1 Sample DAXPY - Numbers are erroneous

Figure 2: Performance of Vector Addition with Scale

## 8.2 Sample DGEMV - Numbers are erroneous

Vendor BLAS performance for DGEMV at CEWES MSRC

Figure 3: Performance of Matrix-Vector Multiplication with Scale